



Sudoku Solver



Ingo Prötzel, Christoph Hanke,
Benjamin Müller



Introduction Sudoku

Fill in the digits 1 - 9 so that

each row and

each column and

each 3x3 sub-grid

contains each digit exactly once.

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9


Data

Four different datasets from Kaggle:

- <https://www.kaggle.com/datasets/bryanpark/sudoku>
1 million puzzles
- <https://www.kaggle.com/datasets/rohanrao/sudoku>
9 million puzzles
- <https://www.kaggle.com/datasets/radcliffe/3-million-sudoku-puzzles-with-ratings>
3 million puzzles
- <https://www.kaggle.com/datasets/informoney/4-million-sudoku-puzzles-easytohard>
4 million puzzles

Data Content

- 17 000 000 solutions are 16 999 295 unique
- between 1 - 64 missing digits



Base Model

- Convolutional Neural Network-based Sudoku Solver
- Sudokus are zero mean-centred normalized ($[-0.5, .05]$) and represented by a $(9, 9, 1)$ -Tensor
- the model transforms every sudoku field into a multi-class classification of size 9
- -> size of probability tensor is 81×9 which is converted to a 9×9 solution tensor using softmax
- Loss function is `sparse_categorical_crossentropy`

Base Model

```
keras.models.Sequential([
    Conv2D(filters=64, kernel_size=(3,3), activation="relu", padding='same', input_shape=(9,9,1)),
    BatchNormalization(),
    Conv2D(filters=64, kernel_size=(3,3), activation="relu", padding='same'),
    BatchNormalization(),
    Conv2D(filters=128, kernel_size=(1,1), activation="relu", padding='same'),

    Flatten(),
    Dense(81*9),
    Reshape((-1, 9)),
    Activation('softmax')
])
```

Base Model

- First approach: predict everything at once
 - -> bad performance
- Second approach: predict only one field per iteration
 - take the value with highest probability and start all over again with $n+1$ given values
 - human centered approach
 - takes very long (~ 20 seconds per sudoku)
 - higher reliability (99% on certain datasets)

Sums and Products

How to improve this already good base model?

- Looking for ways to make it faster and/or even more reliable
- Let's support the model with additional information!
 - every column and row has a sum of 45 ($\sum_{i=1}^9$)
 - every column and row has a product of 362.880 ($\prod_{i=1}^9$)
 - expand the 9x9 tensor to a 10x10 resp. 11x11 tensor
 - Idea: the CNN uses some kind of linear regression which is interwoven to the inference process

[0	0	4	3	0	0	2	0	9	18]
[0	0	5	0	0	9	0	0	1	15]
[0	7	0	0	6	0	0	4	3	20]
[0	0	6	0	0	2	0	8	7	23]
[1	9	0	0	0	7	4	0	0	21]
[0	5	0	0	8	3	0	0	0	16]
[6	0	0	0	0	0	1	0	5	12]
[0	0	3	5	0	8	6	9	0	31]
[0	4	2	9	1	0	3	0	0	19]
[7	25	20	17	15	29	16	21	25	0]

[illegible]

Sums and Products

Difficulties:

- Model becomes quite big:

- `model.add(Dense(81*9))` `#(9x9x9)`
- `-> model.add(Dense(100*45))` `#(10x10x45)`
- `-> model.add(Dense(121*362.880))` `#(11x11x362.880)`

Sums and Products

Difficulties:

- Model becomes quite big:
 - `model.add(Dense(81*9))` `#(9x9x9)`
 - `-> model.add(Dense(100*45))` `#(10x10x45)`
 - `-> model.add(Dense(121*362.880))` `#(11x11x362.880)`
- "45" and "362.880" is always the correct answer
 - > Model doesn't care at all about the additional information!

Custom loss function

- Make loss function aware of Sudoku specifics
- Check the difference in the count of each digit
- Use the difference to modify the given 'regular' loss function
- Process as tensors

Custom loss function

```
def custom_loss_function(y_true, y_pred):  
    error = []  
    y_pred1=tf.math.argmax(y_pred,axis=2)  
    a, b = y_pred1.get_shape()  
    for i in range(a):  
        _, _, cnts = tf.unique_with_counts(y_pred1[i,:])  
        cnts_diff = tf.abs(cnts-9)  
        pred_error = tf.reduce_sum(cnts_diff) / 144  
        error.append(pred_error)  
    corrector = 1 + tf.square(tf.reduce_mean(error))  
    corrector = tf.cast(corrector, tf.float32)  
    scce = tf.keras.losses.SparseCategoricalCrossentropy()  
    result = scce(y_true, y_pred)  
    return result * corrector
```

Reduce runtime - reduce model complexity


```
keras.models.Sequential([
    Conv2D(filters=64, kernel_size=(3,3), activation="relu", padding='same', input_shape=(9,9,1)),
    BatchNormalization(),
    Conv2D(filters=64, kernel_size=(3,3), activation="relu", padding='same'),
    BatchNormalization(),
    Conv2D(filters=128, kernel_size=(1,1), activation="relu", padding='same'),

    Flatten(),
    Dense(81*9),
    Reshape((-1, 9)),
    Activation('softmax')
])
```

1. Delete Layers
2. Reduce number of filters

Reduce runtime - pruning

- remove redundant neurons
- remove rarely activated neurons



Reduce runtime - Results


Model	Learntime	#Param (in Mio)		testtime	<u>iterative</u> complete	relative		testtime	<u>oneshot</u> complete	relative
64, 64, 64, 128	7214	7.6		257.72	1	1		6.32	0	0.70
64, 64, 128	1581	7.6		223.99	1	1		5.11	0	0.72
64, 64, 128 P	2529	7.6		252.42	1	1		6.08	0	0.72
64 ,64	1010	3.8		215.86	1	1		4.82	0	0.70
32, 32	575	1.9		211.36	0.96	0.64		5.10	0	0.67
16, 16	366	0.9		210.96	1	1		5.07	0	0.70
16, 16 P	467	7.6		220.55	1	1		5.35	0	0.70
8, 8	259	0.5		209.18	0	0.47		4.97	0	0.52
64	783	3.8		215.60	0	0.31		4.97	0	0.35

Reduce runtime - Result & followup measures

- model reduction without losing accuracy is possible
- BUT: reducing model complexity had a low effect on testtime
- use better Pruning methods
- time is lost in the interactive process due to array methods?
 - enhance by using faster array methods
- go for the oneshot approach
 - make larger models and prune them
 - Is it possible to reach 100% with this approach?

Picture of first 6 filters

- black means low activation
- white high activation



Lessons learned

- Get a better understanding at reading error messages
- Good looking learning curve doesn't automatically mean good predictions
- Shaping is a pain!
- What is a tensor and what is a numpy array?
- Training takes time
- Batch size matters
- AI is not the solution to everything
- ChatGPT is better in generating a classical backtracking solution than solving sudokus itself

The End

Questions?